

Lecture 2: Bellman operator, Banach's fixed point, solving MDPs

SUMS 707 - Basic Reinforcement Learning

Gabriela Moisescu-Pareja and Viet Nguyen

McGill University, Mila

January 21, 2021


Recap of last lecture

Last time we introduced:

- Markov Decision Processes
- Policies
- State value functions, action-value functions
- Optimality

Recap of last lecture: RL, MDPs I

- Agent is in some state, performs actions
- Environment transitions agent to a state, gives a reward
- This interactive process is *modeled* through a Markov Decision Process (MDP)
- Assume the *Markov property*



Recap of last lecture: RL, MDPs II

- $M = (\mathcal{S}, \mathcal{A}, \mathcal{P}_0)$ is an MDP, where \mathcal{S} are states, \mathcal{A} are actions, and \mathcal{P}_0 is a transition kernel
- the transition kernel gives a reward function $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ and a *state transition kernel* \mathcal{P} :

$$\mathbb{P}(s'|s, a) = \mathcal{P}(s, a, s')$$

- Agent at state s_t takes action a_t , the environment transitions the agent to state s_{t+1} and gives reward r_{t+1} as $(s_{t+1}, r_{t+1}) \sim \mathcal{P}_0(\cdot | s_t, a_t)$
- Want to maximize expected return, i.e. the expectation of

$$\mathcal{R} = \sum_{t=0}^{\infty} \gamma^t r_{t+1}$$

Recap of last lecture: Policies I

- Policies formalize an agent's behavior
- A policy π is a mapping $\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$
- The agent in state s_t samples an action a_t according to $a_t \sim \pi(\cdot | s_t)$
- Deterministic policies: $a_t = \pi(s_t)$

Goal: Find the optimal policy, i.e. the policy that results in the maximum expected returns.

- How? \rightarrow Need to quantify how “good” a state is under some policy

Recap of last lecture: Value functions

- Given some policy, need to quantify the value of a state. Define value functions and action-value functions:

$$V^{\pi}(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1} \middle| s_0 = s, \pi \right], \quad s \in \mathcal{S}$$

$$Q^{\pi}(s, a) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1} \middle| s_0 = s, a_0 = a, \pi \right], \quad s \in \mathcal{S}, a \in \mathcal{A}$$

- Answers the questions “Given that I’m playing policy π , what is the value of being in state s ?” and “Given that I’m playing policy π , what is the value of being in state s and taking action a ?”

Recap of last lecture: Optimality

Define the optimal value and action-value functions:

$$V^*(s) = \max_{\pi} V^{\pi}(s)$$

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a)$$

- Intuitively, they specify the best performance in the MDP
- If we have Q^* , we can play the greedy policy
 $\pi^*(s) = \arg \max_{a \in \mathcal{A}} Q^*(s, a)$
- By playing π^* , you're getting the most rewards possible
- Pog

Plan for today

- Bellman equations, Bellman optimality equations
- Banach fixed point theorem and contractions
- Solving MDPs with DP
 - Policy iteration
 - Value iteration

Bellman Equations I

Let's look at our familiar ravioli:

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1} \mid s_0 = s, \pi \right]$$

There's a pickle here: this involves an infinite computation. Computers don't like infinite computations. Is there anything we can do?

Yes! We unpickle the pickle!

$$\begin{aligned} V^\pi(s) &= \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r_{t+1} \mid s_0 = s, \pi \right] \\ &= \mathbb{E} [r_1 + \gamma r_2 + \gamma^2 r_3 + \dots \mid s_0 = s, \pi] \\ &= \mathbb{E} [r_1 + \gamma V^\pi(s_1) \mid s_0 = s, \pi] \end{aligned}$$

Bellman Equations II


A similar decomposition gives us:

$$Q^\pi(s, a) = \mathbb{E} [r_1 + \gamma Q^\pi(s_1, a_1) | s_0 = s, a_0 = a, \pi]$$

These are things to keep in mind moving forward...

Bellman Equations for V^π I

There seems to be some recursion going on...




- In state s , we can take actions $a \in \mathcal{A}$ with probability $\pi(a|s)$
- By definition, we have that

$$\begin{aligned} V^\pi(s) &= \mathbb{E}_{a \sim \pi(\cdot|s)} [Q^\pi(s, a)] \\ &= \sum_{a \in \mathcal{A}} \pi(a|s) Q^\pi(s, a) \end{aligned}$$

Bellman Equations for V^π II

One step look-ahead gives us:



- The environment moves you to state s' with probability $\mathbb{P}(s'|s, a)$, gives you reward r .

$$\begin{aligned} V^\pi(s) &= \sum_{a \in \mathcal{A}} \pi(a|s) Q^\pi(s, a) \\ &= \sum_{a \in \mathcal{A}} \pi(a|s) \left(r_{(s,a)} + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) V^\pi(s') \right) \end{aligned}$$


Deterministic Bellman Equations for V^π

When π is deterministic, this equation reduces to:

$$V^\pi(s) = r(s, \pi(s)) + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) V^\pi(s')$$

where we used here the reward function $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ defined last time!


Bellman Equations for Q^π I



- Taking action a in state s , we obtain a reward $r_{(s,a)}$ and transitions to s' w.p. $\mathbb{P}(s'|s, a)$
- Thus, we have that

$$Q^\pi(s, a) = r_{(s,a)} + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) V^\pi(s')$$

Bellman Equations for Q^π II



- In state s' , we choose action a' according to π . But we already have a formula for V^π from the previous slides!

$$\begin{aligned} Q^\pi(s, a) &= r_{(s,a)} + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) V^\pi(s') \\ &= r_{(s,a)} + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) \mathbb{E}_{a' \sim \pi(\cdot|s)} [Q^\pi(s', a')] \\ &= r_{(s,a)} + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) \sum_{a' \in \mathcal{A}} \pi(a'|s') Q^\pi(s', a') \end{aligned}$$

Deterministic Bellman Equations for Q^π

$$Q^\pi(s, a) = r_{(s,a)} + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) \sum_{a' \in \mathcal{A}} \pi(a'|s') Q^\pi(s', a')$$

When π is deterministic, this equation reduces to

$$Q^\pi(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) Q^\pi(s', \pi(s'))$$

Bellman Operator I

Let's stare again real hard at the two *deterministic* equations we just wrote down:

$$V^\pi(s) = r(s, \pi(s)) + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) V^\pi(s')$$

$$Q^\pi(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) Q^\pi(s', \pi(s'))$$

What do you see?

Bellman Operator II

$$V^\pi(s) = r(s, \pi(s)) + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) V^\pi(s')$$

$$Q^\pi(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) Q^\pi(s', \pi(s'))$$

Yes! It does look like we can define an *operator* underlying π . Define the Bellman operator $T^\pi : (\mathcal{S} \rightarrow \mathbb{R}) \rightarrow (\mathcal{S} \rightarrow \mathbb{R})$ defined by its actions on \mathcal{S} via *any* $V : \mathcal{S} \rightarrow \mathbb{R}$ in the following way:

$$(T^\pi V)(s) = r(s, \pi(s)) + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) V(s')$$

Bellman Operator III

In the same way, one can define the Bellman operator for functions of $\mathcal{S} \times \mathcal{A}$. Consider $T^\pi : (\mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}) \rightarrow (\mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R})$ defined by:

$$(T^\pi Q)(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) Q(s', \pi(s'))$$

Bellman Operator IV

In this way, we can rewrite the Bellman equations as:

$$T^\pi V^\pi = V^\pi, \quad T^\pi Q^\pi = Q^\pi$$

Why do we care about this?

- This system is *linear*, T^π is an affine linear operator.
- When $\gamma \in (0, 1)$, T^π is a max-norm *contraction* and the fixed-point equation $T^\pi V = V$ has an *unique* solution.
- The unique solution is *exactly* V^π !
- Similarly for Q^π ...

Optimal Policy

Define a partial ordering over policies:

$$\pi \geq \pi' \text{ if } \forall s \in \mathcal{S}, V^\pi(s) \geq V^{\pi'}(s)$$

Theorem

For any MDP, there is always an optimal policy π^ (not necessarily unique). All optimal policies achieve the optimal value function $V^{\pi^*}(s) = V^*(s)$ and optimal action-value function $Q^{\pi^*}(s, a) = Q^*(s, a)$.*

There will always be a deterministic optimal policy, so if we know Q^* we already have the optimal policy by playing greedily.

Optimality, continued

Recall the definitions of the optimal value and action-value functions:

$$V^*(s) = \max_{\pi} V^{\pi}(s), \quad Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a)$$

As we stated previously, they can actually be rewritten in terms of each other:

$$V^*(s) = \sup_a Q^*(s, a)$$

and

$$Q^*(s, a) = r_{(s,a)} + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) V^*(s')$$

Exercise: Optimality equations

Exercise: We previously derived the Bellman equations for V^π and Q^π . The idea now is to derive the Bellman equations for V^* and Q^* (we call these the Bellman optimality equations) by following the exact same process. This time, however, we know that $V^* = V^{\pi^*}$ where π^* is an optimal deterministic policy.

Try to rewrite V^* and Q^* in terms of each other when you draw out the trees (like we did).

Exercise: Hint

You should arrive at something that looks like this:

$$V^*(s) = \sup_a \left\{ R_{(s,a)} + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) V^*(s') \right\}$$

$$Q^*(s, a) = r_{(s,a)} + \gamma \sum_{s' \in \mathcal{S}} \mathbb{P}(s'|s, a) \sup_{a'} Q^*(s', a')$$

Bellman optimality equations, why?

- Many decision-making methods can be viewed as ways of approximately solving the Bellman optimality equations
- DP methods like policy and value iterations are very closely related to the Bellman optimality equation
- In the RL setting where we don't know the underlying environment dynamics, many methods to finding a good policy can be understood as approximately solving the Bellman optimality equations, but using actual experienced transitions instead of the knowledge of the expected transitions
- We will look into this more closely in the near future

Solving MDPs I

We are now ready to solve MDPs! There are various methods that exist to do this, and we list some of the fundamental methods:

- **Dynamic programming** (which we will focus on today)
- Monte Carlo methods
- Temporal-difference methods

Solving MDPs II


Of course, each method has its strengths and weaknesses.

- Dynamic programming:
 - Well-developed mathematically (in fact what we introduced in the previous section will come in handy now)
 - However, it requires a **full** description of the model of the environment (dynamics and rewards)
- Monte Carlo methods:
 - Don't require full model and are conceptually simple
 - However, they are not especially suited for step-by-step incremental computation
- Temporal-difference methods:
 - Don't require a full model
 - However, they are more complex to analyze (though this might have changed...).

Solving MDPs III

Finally, the methods of course differ in terms of their computational efficiency and speed of convergence.

Solving MDPs using dynamic programming I



A fun xkcd comic.

Solving MDPs using dynamic programming II

We proceed by taking the dynamic programming (DP) approach to solving an MDP. The problem of finding π^* given full specification of the MDP (model of the environment) is generally known as *planning*. There are two classical methods that we will go over:

- Policy iteration
- Value iteration

Quick Dynamic Programming Overview I

DP allows us to solve complex problems by breaking them down into subproblems, solving the subproblems, and combining the solutions to those subproblems to arrive at a solution for the original problem. There are two properties:

- Optimal substructure
 - Principle of optimality applies
 - Optimal solution can be decomposed into subproblems
- Overlapping subproblems
 - Subproblems recur many times
 - Solutions can be cached and reused

Quick Dynamic Programming Overview II

As it turns out, MDPs satisfy this!

- Bellman equation gives recursive decomposition
- The value function stores and reuses solutions

Quick Dynamic Programming Overview III

As we will see, these DP algorithms will be obtained by turning Bellman equations into update rules for improving approximations of the desired value functions.

Policy evaluation I

Given a policy π , we would like to compute V^π . This is called *policy evaluation* (and sometimes called the prediction problem).

Recall

$$\begin{aligned} V^\pi &= \mathbb{E} \left[r_{t+1} + \gamma V^\pi(s_{t+1}) \mid s_t = s, \pi \right] \\ &= \sum_a \pi(s, a) \sum_{s'} \mathbb{P}(s' | s, a) [r(s, a) + \gamma V^\pi(s')] \end{aligned}$$

Policy evaluation III

Computing V^π can be done iteratively and is known as *Iterative Policy Evaluation*. Consider a sequence of approximations V_0, V_1, V_2, \dots


- Initial approximation V_0 is chosen arbitrarily (though often set to 0)
- Each successive approximation is obtained by using the Bellman equation for V^π as an update rule:

$$\begin{aligned} V_{k+1}(s) &= \mathbb{E} \left[r_{t+1} + \gamma V_k(s_{t+1}) \mid s_t = s, \pi \right] \\ &= \sum_a \pi(s, a) \sum_{s'} \mathbb{P}(s' | s, a) [r(s, a) + \gamma V_k(s')] \end{aligned}$$

for all states $s \in S$.

We have that $V_k = V^\pi$ is a fixed point for this update rule thanks to the Bellman equation for V^π . The sequence $\{V_k\}$ can be shown, in general, to converge to V^π as $k \rightarrow \infty$.

Policy evaluation IV



We can understand the procedure graphically.

Policy improvement I

Just now, we computed the value function given a policy π . This turns out to be useful to help us find *better policies*!

Policy improvement II

Suppose we have computed the value function V^π for a deterministic policy π .

- For some state s , we might like to know whether we should actually change the policy to deterministically choose another action $a \neq \pi(s)$.
- We know how good it is to follow the current policy from s thanks to V^π .
- But we might wonder if it is better or worse to change to the new policy where we choose action a instead of action $\pi(s)$ at state s .

Policy improvement III

One way to answer this question is to consider choosing a when in s and following the existing policy π thereafter. We can see how good this is by looking at the state-action value function:

$$\begin{aligned} Q^\pi(s, a) &= \mathbb{E} \left[r_{t+1} + \gamma V^\pi(s_{t+1}) \mid s_t = s, a_t = a, \pi \right] \\ &= \sum_{s'} \mathbb{P}(s' | s, a) [r(s, a) + \gamma V^\pi(s')] \end{aligned}$$

Is this greater than or less than V^π ? If it is greater, it would be better to select a and then follow π than to follow π the entire time. This key idea relies on the following *Policy Improvement Theorem*

Policy improvement IV

Theorem (Policy Improvement Theorem)

Choose some stationary policy π_0 and let π be greedy w.r.t.

$V^{\pi_0} : T^{\pi} V^{\pi_0} = T^ V^{\pi_0}$. Then $V^{\pi} \geq V^{\pi_0}$, i.e. π is an improvement upon π_0 .*

In particular, if $T^ V^{\pi_0}(s) > V^{\pi_0}(s)$ for some state s , then π strictly improves upon π_0 at s , so*

$$V^{\pi}(s) > V^{\pi_0}(s).$$

On the other hand, when $T^ V^{\pi_0} = V^{\pi_0}$ then π_0 is an optimal policy.*

The proof is omitted for now, but we will get to it next class.

Policy iteration I

Policy iteration is an algorithm for the *control problem*. At a high-level, policy iteration performs the following steps:

- 1 Policy evaluation: value function is computed for the currently policy
- 2 Policy improvement: the policy is made greedy w.r.t. the value function
- 3 Repeat steps 1, 2 until convergence to optimal policy.

Policy iteration II

Here we offer a run-down:

- We start with a deterministic policy π , where at state s we act according to $a = \pi(s)$
- We can improve the policy by acting greedily

$$\pi'(s) = \arg \max_{a \in A} Q^\pi(s, a)$$

- This improves the value from any state s over one step

$$Q^\pi(s, \pi'(s)) = \max_{a \in A} Q^\pi(s, a) \geq Q^\pi(s, \pi(s)) = V^\pi(s)$$

Policy iteration III

- This improves the value function, $V^{\pi'}(s) \geq V^{\pi}(s)$

$$\begin{aligned} V^{\pi}(s) &\leq Q^{\pi}(s, \pi'(s)) \\ &= \mathbb{E} \left[r_{t+1} + \gamma V^{\pi}(s_{t+1}) \middle| s_t = s, \pi' \right] \\ &\leq \mathbb{E} \left[r_{t+1} + \gamma Q^{\pi}(s_{t+1}, \pi'(s_{t+1})) \middle| s_t = s, \pi' \right] \\ &\leq \mathbb{E} \left[r_{t+1} + \gamma r_{t+2} + \gamma^2 Q^{\pi}(s_{t+2}, \pi'(s_{t+2})) \middle| s_t = s, \pi' \right] \\ &\leq \mathbb{E} \left[r_{t+1} + \gamma r_{t+2} + \dots \middle| s_t = s \right] \\ &= V^{\pi'}(s) \end{aligned}$$

Policy iteration IV

- If improvements stop,


$$Q^{\pi}(s, \pi'(s)) = \max_{a \in A} Q^{\pi}(s, a) = Q^{\pi}(s, \pi(s)) = V^{\pi}(s)$$

- Then the Bellman optimality equation has been satisfied

$$V^{\pi}(s) = \max_{a \in A} Q^{\pi}(s, a)$$

Therefore $V^{\pi}(s) = V^*(s)$ for all states s and π is an optimal policy.

Policy iteration V



Value iteration I

Idea: Optimize value function directly and then induce a policy. This means that unlike in policy iteration, we don't need to go back and forth between value functions and policies.

Value iteration II

We want to compute V^* directly, and it can be done through an iterative method known as *Value Iteration*. Consider a sequence of approximations V_0, V_1, V_2, \dots

- Initial approximation V_0 is chosen arbitrarily.
- Each successive approximation is obtained by

$$\begin{aligned} V_{k+1}(s) &= \max_{a \in A} \mathbb{E} [r_{t+1} + \gamma V_k(s_{t+1}) | s_t = s, a_t = a] \\ &= \max_{a \in A} \sum_{s'} \mathbb{P}(s' | s, a) [r(s, a) + \gamma V_k(s')] \end{aligned}$$

for all $s \in S$.


And $\{V_k\}$ can be shown to converge to V^* under the same conditions that guarantee the existence of V^* .

Value iteration III

On the convergence:

- We have that the Bellman optimality operator T^* has a unique fixed point
- V^* is a fixed point of T^* by the Bellman optimality equation
- By the Banach fixed point theorem, value iteration converges to V^* at a geometric rate.

Value iteration IV



We can understand the procedure graphically.

Important thingies

- MDPs are perfectly suited to DP given by the Bellman equations.
- DP methods such as PI and VI require a full model and are especially good for finite MDPs of small to moderate size, but other methods are necessary to tackle more complex RL problems
- Convergence of the PI and VI follow from the Banach fixed point theorem.

Next lecture

What do we do when we don't know the dynamics of the environment?

- Model-free prediction and control
- TD and MC methods, Q-learning
- (maybe) Policy improvement theorem proof

References

- Reinforcement Learning: Theory and Algorithms by Alekh Agarwal, Nan Jiang, Sham M. Kakade
- Algorithms for Reinforcement Learning by Csaba Szepesvári
- Reinforcement Learning: An Introduction by Andrew Barto and Richard S. Sutton
- “Introduction to Reinforcement Learning” Lectures by David Silver
- “CS 598 - Statistical Reinforcement Learning” Notes by Nan Jiang